

DeNovoSync: Efficient Support for Arbitrary Synchronization without Writer-Initiated Invalidations



Hyojin Sung and Sarita Adve

**Department of Computer Science
University of Illinois, EPFL**

Motivation

Complex software

Data races, non-determinism, implicit communication, ...

Shared Memory

Complex, inefficient hardware

Complex coherence, consistency, unnecessary traffic, ...

Motivation

Complex software

Data races, non-determinism, implicit communication, ...

WILD Shared Memory

Complex, inefficient hardware

Complex coherence, consistency, unnecessary traffic, ...

Motivation

Complex software

Data races, non-determinism, implicit communication, ...

Disciplined
~~WILD~~ Shared Memory

Complex, inefficient hardware

Complex coherence, consistency, unnecessary traffic, ...

Motivation

Complex software

Data races, non-determinism, implicit communication, ...

**Structured synch +
Explicit memory
side effects**

Disciplined Shared Memory


Complex, inefficient hardware

Complex coherence, consistency, unnecessary traffic, ...

Motivation

Simpler
~~**Complex**~~ software

No data races, safe non-determinism, explicit sharing, ...



Structured synch +
Explicit memory
side effects

Disciplined Shared Memory

Complex, inefficient hardware

Complex coherence, consistency, unnecessary traffic, ...

Motivation

Simpler
~~**Complex**~~ software

No data races, safe non-determinism, explicit sharing, ...

Structured synch +
Explicit memory
side effects

Disciplined Shared Memory

Complex, inefficient hardware

Complex coherence, consistency, unnecessary traffic, ...

Motivation

Simpler
~~Complex~~ software

No data races, safe non-determinism, explicit sharing, ...

Structured synch +
Explicit memory
side effects

Disciplined Shared Memory

Simpler, more efficient
~~Complex, inefficient~~ hardware

DeNovo [PACT11], DeNovoND [ASPLOS13, Top Picks 14]

BUT focus on data accesses, **synchronization** restricted

BUT much software (runtime, OS, ...) uses unstructured synch

Motivation



Simpler
~~Complex~~ software

No data races, safe non-determinism, explicit sharing, ...

DeNovoSync:
Support arbitrary synchronization with advantages of DeNovo



DeNovo [PACT11], DeNovoND [ASPLOS13, Top Picks 14]


BUT focus on data accesses, **synchronization** restricted

BUT much software (runtime, OS, ...) uses unstructured synch

Supporting Arbitrary Synchronization: The Challenge

- MESI: Writer sends invalidations to cached copies to avoid stale data

~~⇒ Directory storage, inv/ack msgs, transient states, ...~~



BUT Synchronization?

Naïve: Don't cache synchron

- Prior DeNovo assumptions

~~— Race freedom —~~

~~— Restricted synchronization with special hardware —~~

⇒ Reader self-invalidates stale data

Contributions of DeNovoSync

- **DeNovoSync: Cache arbitrary synch w/o writer invalidations**
- **Simplicity, perf, energy advantages of DeNovo w/o sw restrictions**
- **DeNovoSync vs. MESI for 24 kernels (16 & 64 cores), 13 apps**
 - **Kernels: 22% lower exec time, 58% lower traffic for 44 of 48 cases**
 - **Apps: 4% lower exec time, 24% lower traffic for 12 of 13 cases**


Outline

- **Motivation**
- **Background: DeNovo Coherence for Data**
- **DeNovoSync Design**
- **Experiments**
- **Conclusions**

DeNovo Coherence for Data (1 of 2)

- Original DeNovo software assumptions [PACT'11]
 - Data-race-free
 - Synchronization: Barriers demarcate parallel phases
 - Writeable data regions in parallel phase are explicit **W**
- Coherence
 - Read hit: Don't return stale data
 - Before next parallel phase, cache selectively **self-invalidates** **W**
 - Needn't invalidate data it accessed in previous phase
 - Read miss: Find *one* up-to-date copy
 - Write miss registers at ~~“directory”~~ **registry**
 - Shared LLC data arrays double as registry
 - Keep valid data or registered core id

DeNovo Coherence for Data (2 of 2)



But how to handle arbitrary synchronization?

- DeNovoND adds structured locks [ASPLOS'13, Top Picks'14]
 - When to self-invalidate: at lock acquire
 - What data to self-invalidate: dynamically collected modified data signatures
 - Special hardware support for locks

Outline

- **Motivation**
- **Background: DeNovo Coherence for Data**
- **DeNovoSync Design**
- **Experiments**
- **Conclusions**

Unstructured Synchronization

Michael-Scott non-blocking queue

New node
to be inserted

```
void queue.enqueue(value v):  
  node *w := new node(v, null)  
  ptr t, n  
  loop  
    t := tail  
    n := t->next  
    if t == tail  
      if n == null  
        if (CAS(&t->next, n, w)) break;  
      else CAS(&tail, t, n)  
    CAS(&tail, t, w)
```


- **Data** accesses ordered by synchronization
 - Self-invalidate at synch using static regions or dynamic signatures
- But what about **synchronization**?

DeNovoSync: Software Requirements

- Software requirement: **Data-race-free**
 - Distinguish synchronization vs. data accesses to hardware
 - Obeyed by C++, C, Java, ...
- Semantics: Sequential consistency
- Optional software information for data consistency performance

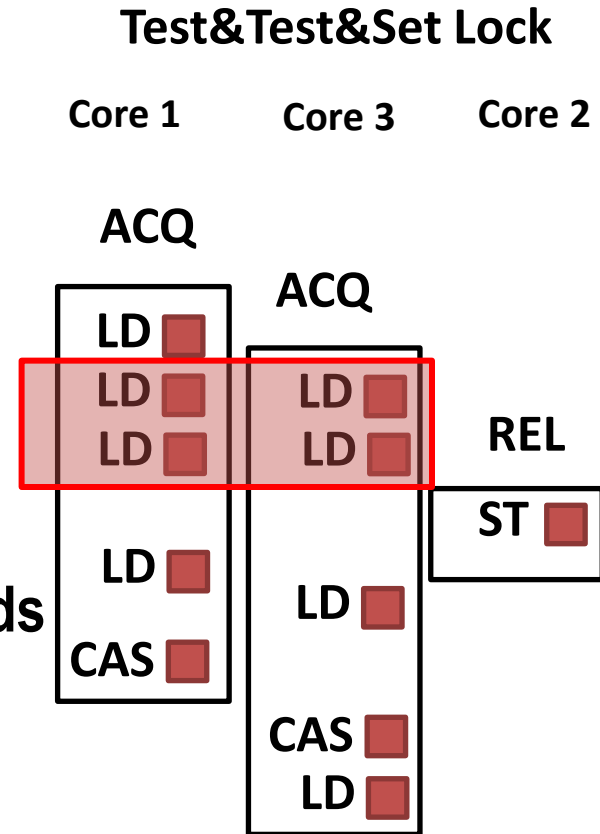
DeNovoSync0 Protocol

- Key: Synch read should not return stale data
- When to self-invalidate synch location?
 - ~~Every synch read?~~
 - Every synch read to non-registered state
- DeNovoSync0 registers (serializes) synch reads
 - Successive reads hit
 - Updates propagate to readers



DeNovoSync0 Protocol

- Key: Synch read should not return stale data
- When to self-invalidate synch location?
 - ~~Every synch read?~~
 - Every synch read to non-registered state
- DeNovoSync0 registers (serializes) synch reads
 - Successive reads hit
 - Updates propagate to readers
 - BUT many registration transfers for Read-Read races



DeNovoSync = DeNovoSync0 + Hardware Backoff

- **Hardware backoff** to reduce Read-Read races
 - Remote synch read requests = hint for contention
 - Delay next (local) synch read miss for backoff cycles

- Two-level **adaptive counters** for backoff cycles

- E = Exponential backoff counter

Read-Read races \Rightarrow Contention \Rightarrow Backoff!

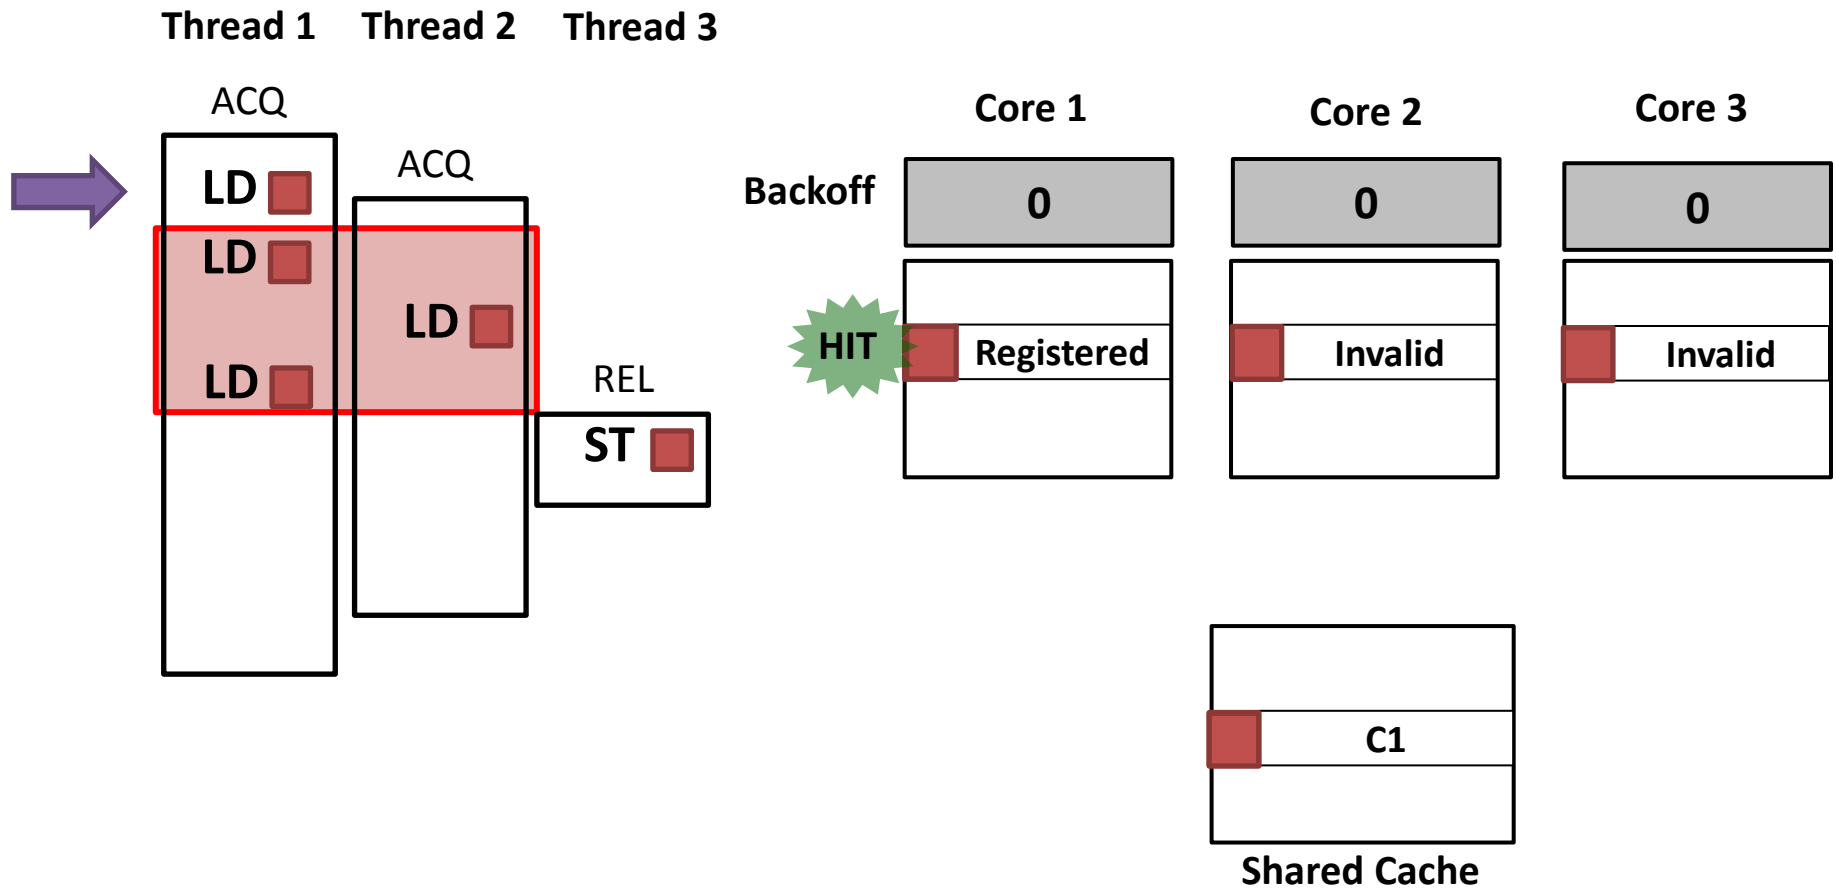
- I = Per-core increment counter

More Read-Read races \Rightarrow More contention \Rightarrow Backoff longer!

- N determined by system configuration

Example

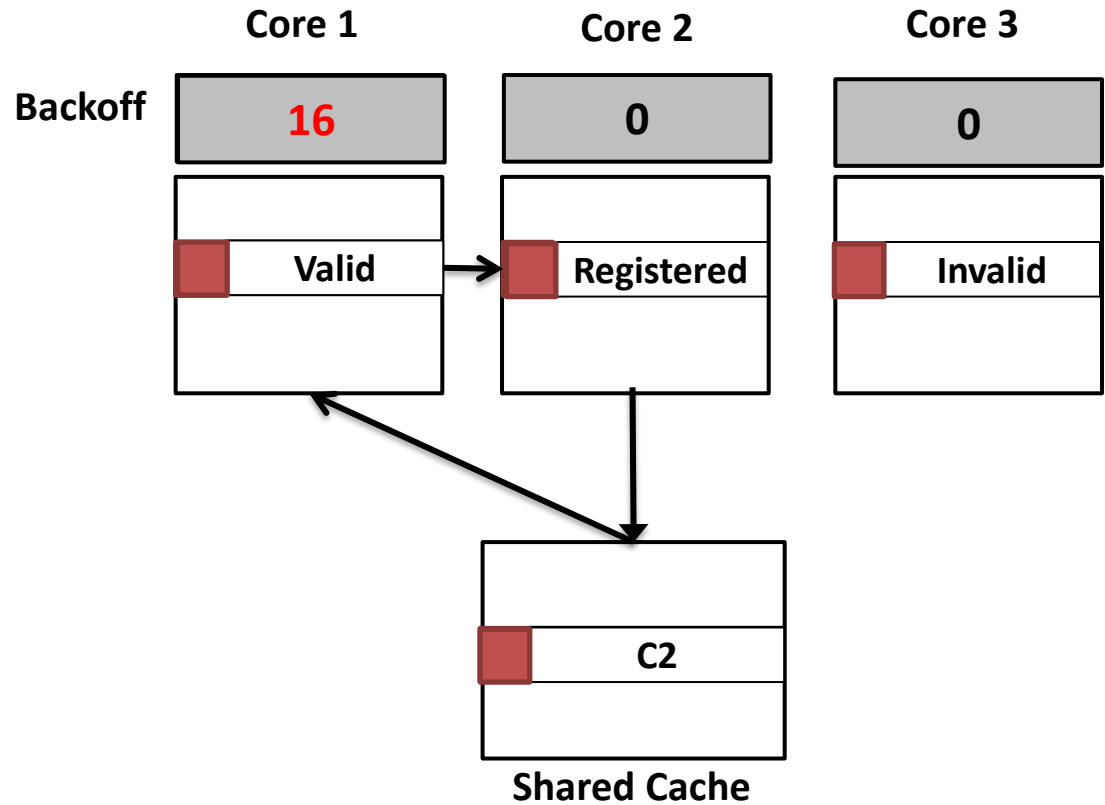

Test&Test&Set Lock



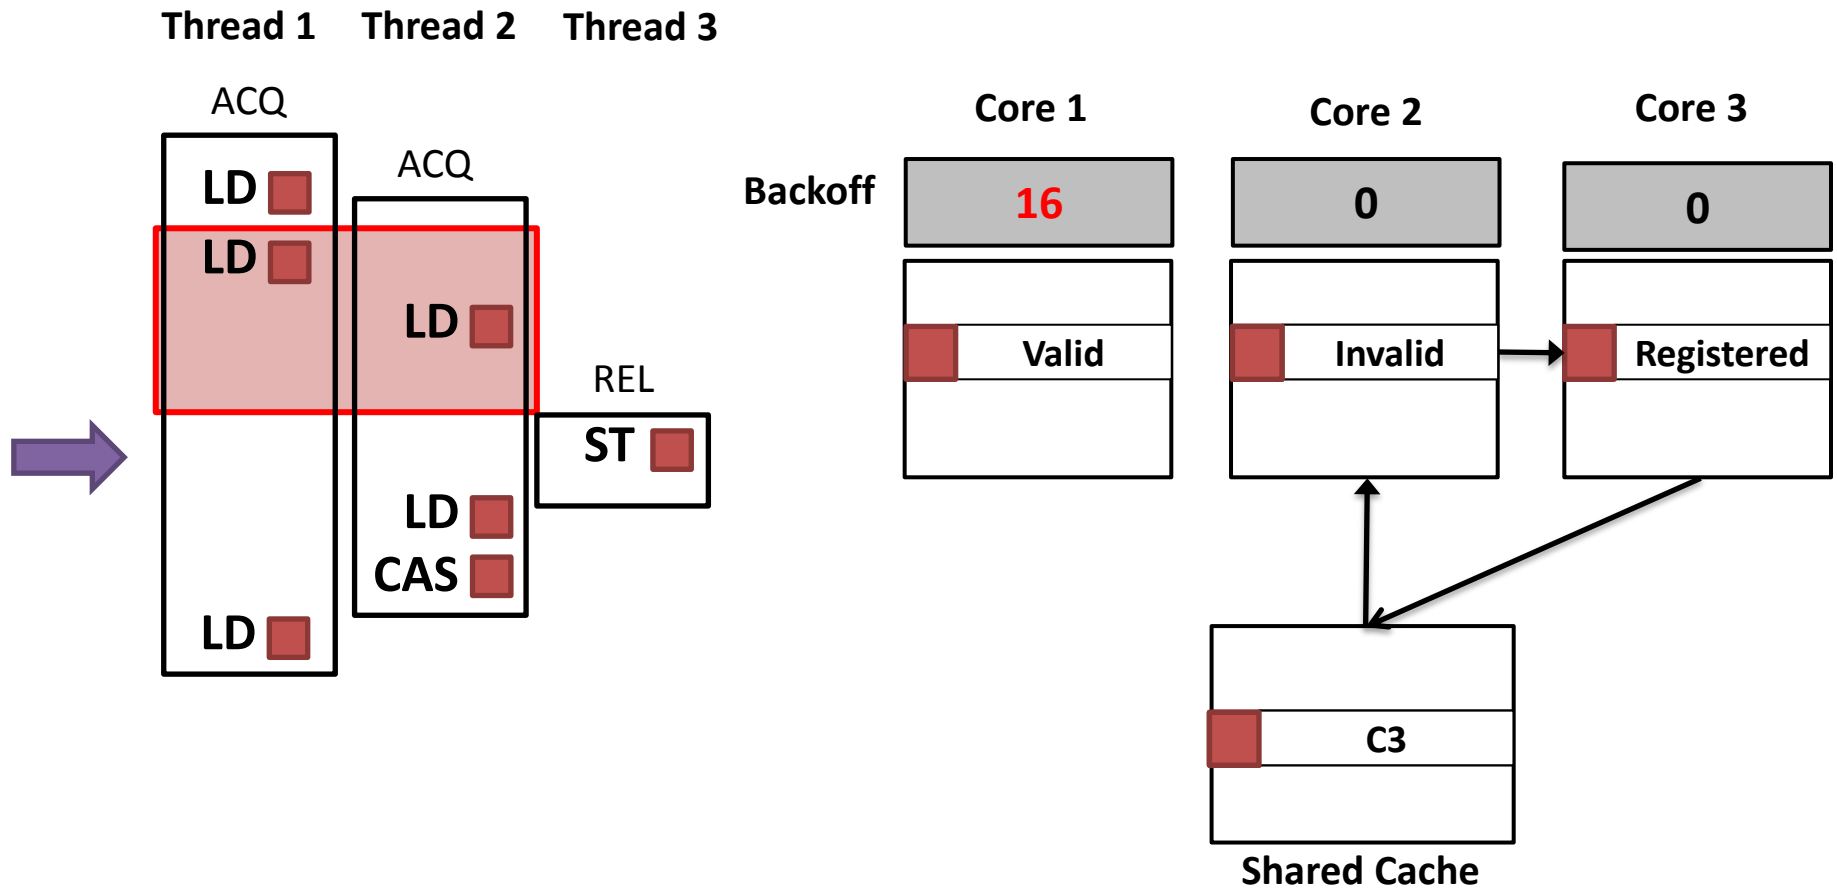
Example

Test&Test&Set Lock

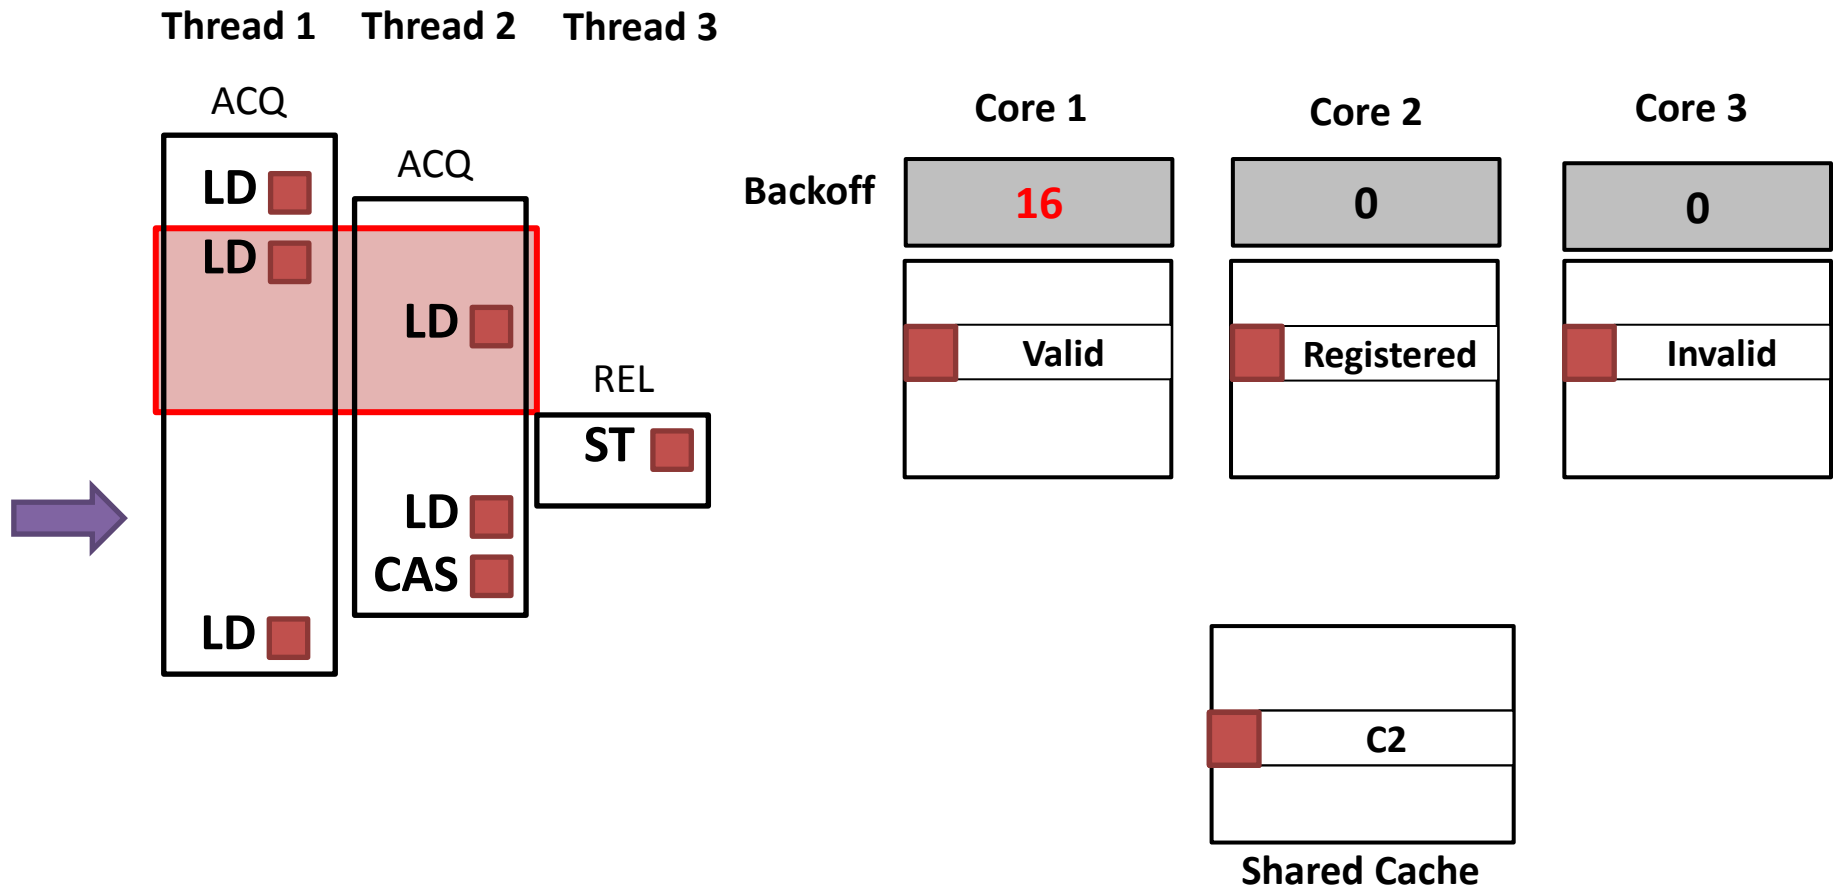
Thread 1 Thread 2 Thread 3



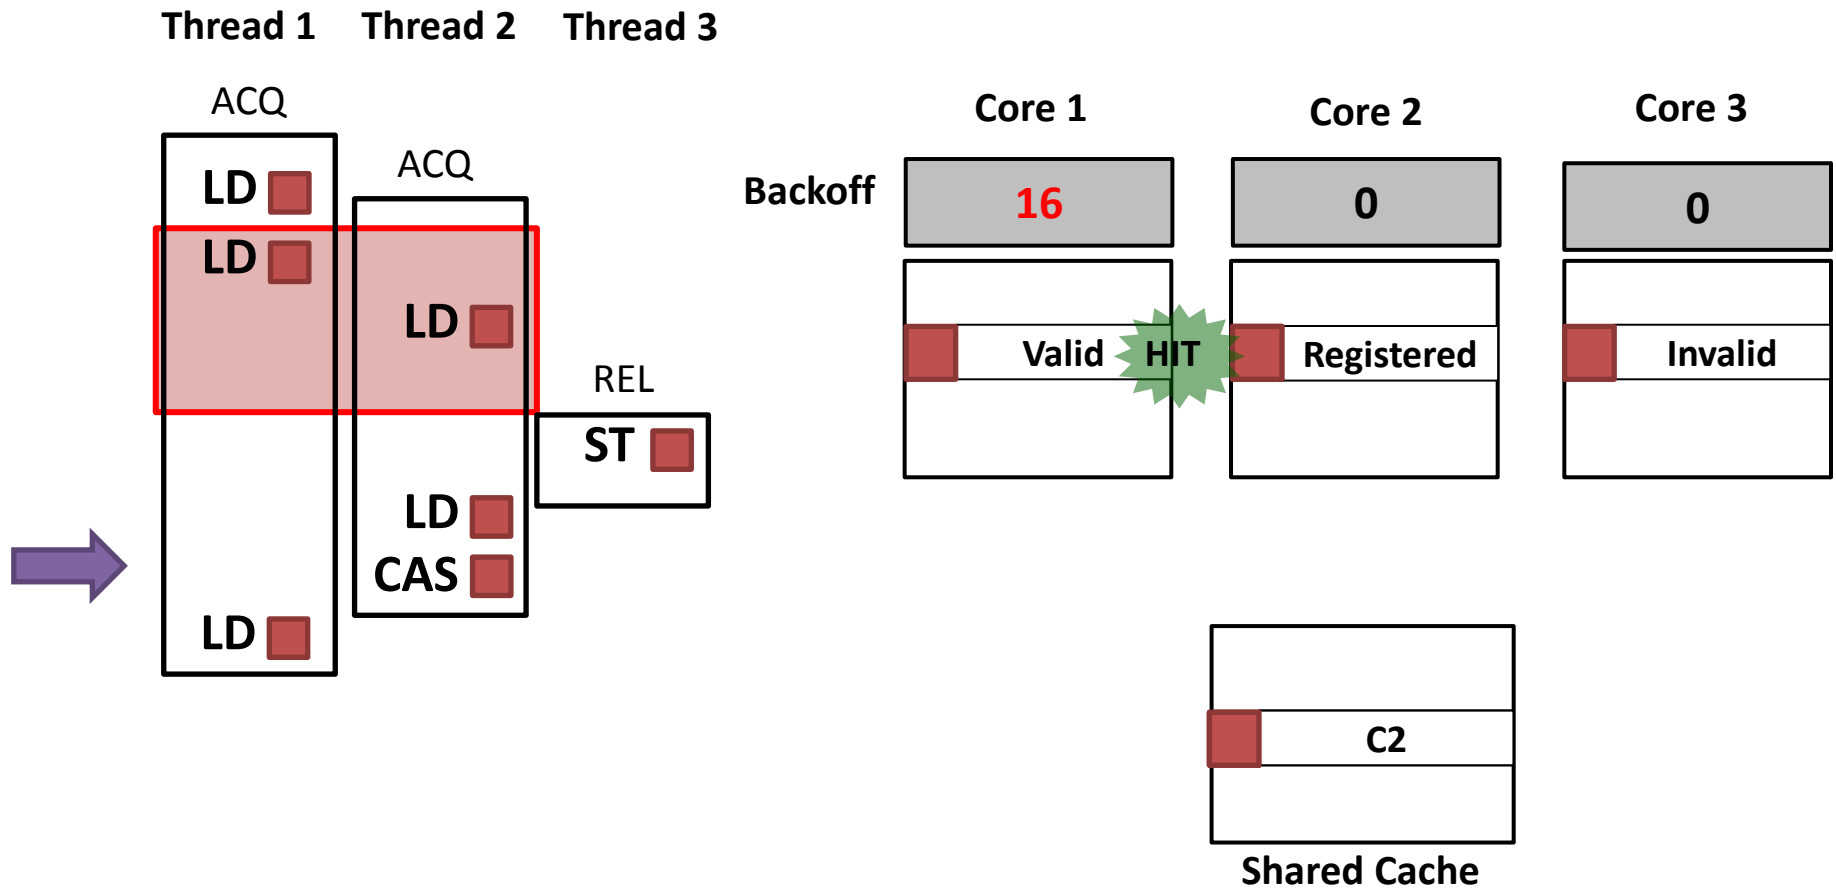
Example



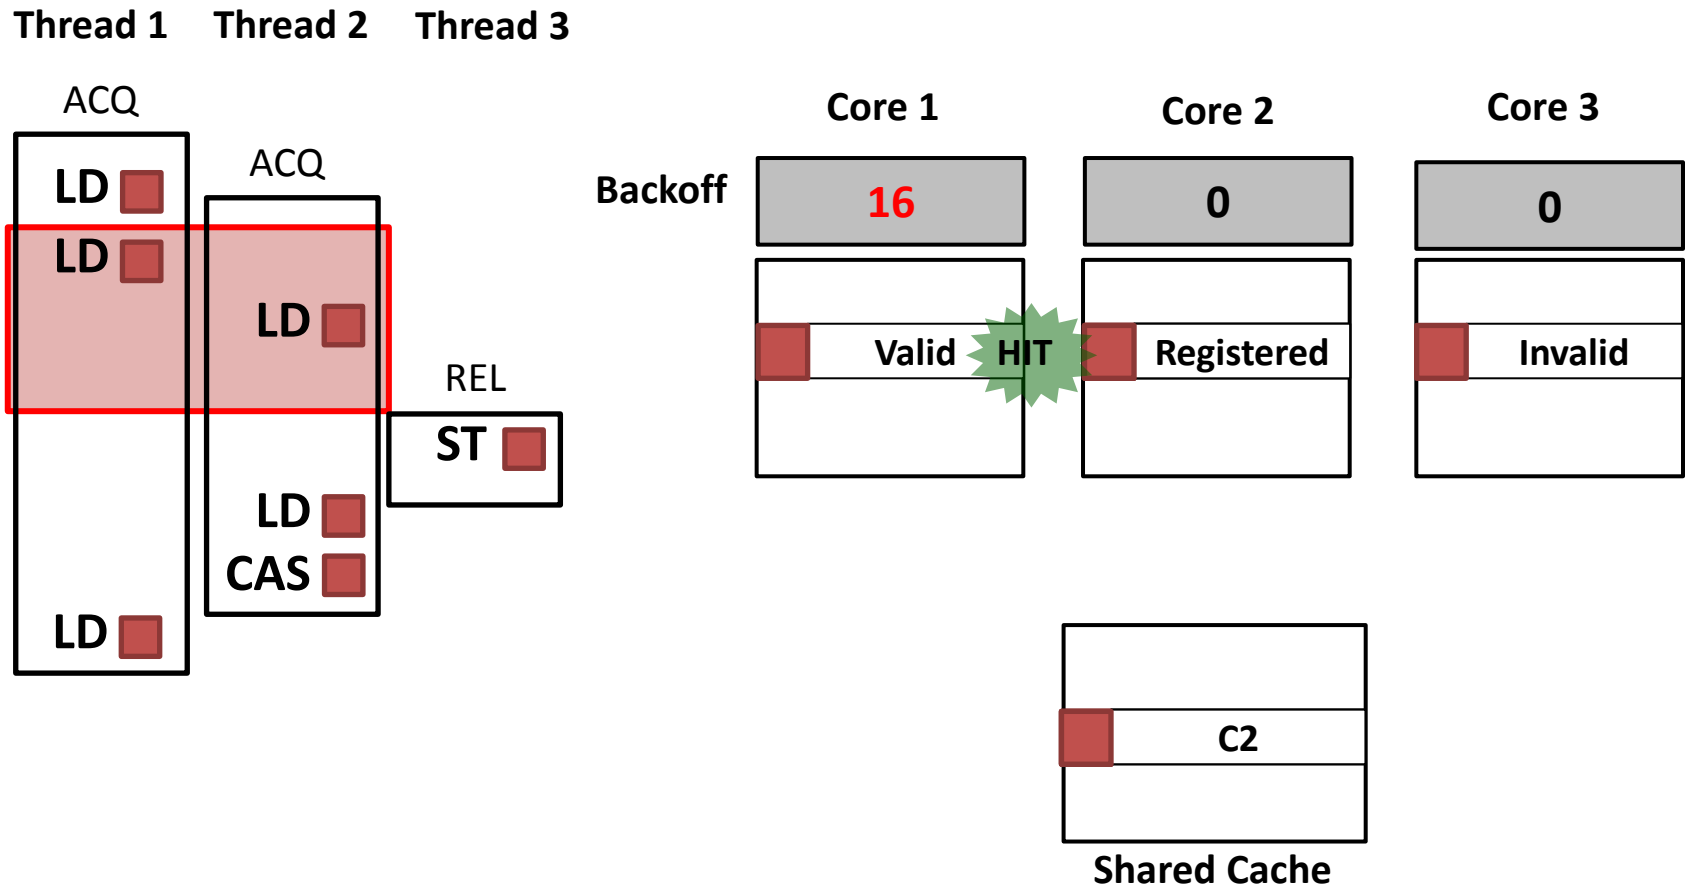
Example



Example



Example



Hardware backoff reduces cache misses from Read-Read races

Outline


- **Motivation**
- **Background: DeNovo Coherence for Data**
- **DeNovoSync Design**
- **Experiments**
 - **Methodology**
 - **Qualitative Analysis**
 - **Results**
- **Conclusions**

Methodology

- Compared MESI vs. DeNovoSync0 vs. DeNovoSync
- Simics full-system simulator
 - GEMS and Garnet for memory and network simulation
- 16 and 64 cores (in-order)
- Metrics: Execution time, network traffic
- Workloads
 - 24 synchronization kernels
 - Lock-based: Test&Test&Set and array locks
 - Non-blocking data structures
 - Barriers: centralized and tree barriers, balanced and unbalanced
 - 13 application benchmarks
 - From SPLASH-2 and PARSEC 3.1
- Annotated data sharing statically (choice orthogonal to this paper)

Qualitative Analysis

- Analyze costs (execution time, traffic) in two parts
 - **Linearization point**
 - Ordering of linearization instruction = ordering of method
 - Usually on critical path
 - **Pre-linearization points**
 - Non-linearization instructions (do not determine ordering)
 - Usually checks, not on critical path



Qualitative Analysis Example (1 of 2)

- Multiple readers, one succeeds: Test&Test&Set locks

Linearization

Pre-linearization

MESI

DeNovoSync0

DeNovoSync

Qualitative Analysis Example (1 of 2)

- Multiple readers, one succeeds: Test&Test&Set locks

	Linearization	Pre-linearization
MESI	Release has high inv overhead, on critical path to next acquire	
DeNovoSync0		
DeNovoSync		

Qualitative Analysis Example (1 of 2)

- Multiple readers, one succeeds: Test&Test&Set locks

	Linearization	Pre-linearization
MESI	Release has high inv overhead, on critical path to next acquire	
DeNovoSync0	No inv overhead	
DeNovoSync		

Qualitative Analysis Example (1 of 2)

- Multiple readers, one succeeds: Test&Test&Set locks

	Linearization	Pre-linearization
MESI	Release has high inv overhead, on critical path to next acquire	
DeNovoSync0	No inv overhead	
DeNovoSync	No inv overhead	

Qualitative Analysis Example (1 of 2)

- Multiple readers, one succeeds: Test&Test&Set locks

	Linearization	Pre-linearization
MESI	Release has high inv overhead, on critical path to next acquire	Local spinning
DeNovoSync0	No inv overhead	
DeNovoSync	No inv overhead	

Qualitative Analysis Example (1 of 2)

- Multiple readers, one succeeds: Test&Test&Set locks

	Linearization	Pre-linearization
MESI	Release has high inv overhead, on critical path to next acquire	Local spinning
DeNovoSync0	No inv overhead	Read-Read races, but not on critical path
DeNovoSync	No inv overhead	

Qualitative Analysis Example (1 of 2)

- Multiple readers, one succeeds: Test&Test&Set locks

	Linearization	Pre-linearization
MESI	Release has high inv overhead, on critical path to next acquire	Local spinning
DeNovoSync0	No inv overhead	Read-Read races, but not on critical path
DeNovoSync	No inv overhead	Backoff mitigates Read-Read races

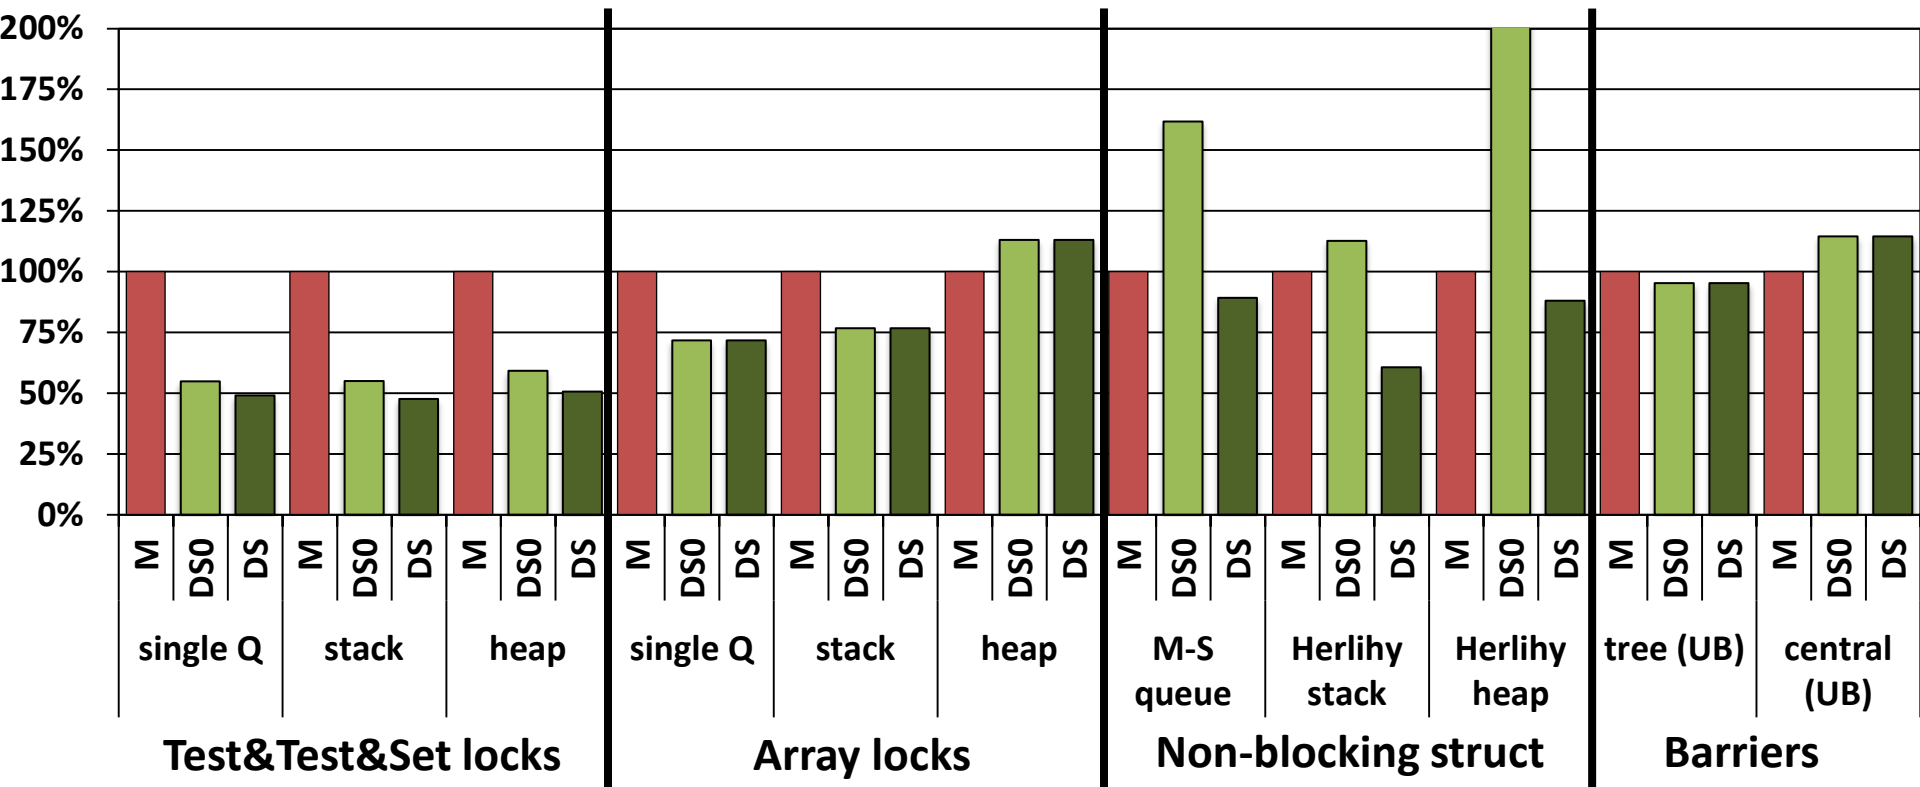
DeNovo expected to be better than MESI

Similar analysis holds for non-blocking constructs

Qualitative Analysis Example (2 of 2)

- **Many readers, all succeed: Centralized barriers**
 - MESI: high linearization due to invalidations
 - **DeNovo: high linearization due to serialized read registrations**
- **One writer, one reader: Tree barriers, array locks**
 - DeNovo, MESI comparable to first order
- **Qualitative analysis only considers synchronization**
 - Data effects: Self-invalidation, coherence granularity, ...
 - Orthogonal to this work, but affect experimental results

Synchronization Kernels: Execution Time (64 cores)



For 44 of 48 cases, 22% lower exec time, 58% lower traffic (not shown)


Remaining 4 cases:

Centralized unbalanced barriers: But tree barriers better for MESI too


Heap with array locks: Need dynamic data signatures for self-invalidation

Applications (64 cores)

Execution Time



Network Traffic



For 12 of 13 cases, 4% lower exec time, 24% lower traffic

Memory time dominated by data (vs. sync) accesses

Conclusions

- **DeNovoSync: First to cache arbitrary synch w/o writer-initiated inv**
 - Registered reads + hardware backoff
 - **With simplicity, performance, energy advantage**
 - No transient states, no directory storage, no invalidation
 - **DeNovoSync vs. MESI**
 - Kernels: For 44 of 48 cases, 22% lower exec time
 - Apps: For 12 of 13 cases, 4% lower exec time
- ⇒ **Complexity-, performance-, energy-efficiency w/o s/w restrictions**
- **Future: DeNovo w/ heterogeneity [ISCA'15], dynamic data signatures**

